

# Package ‘MultFourier’

March 18, 2025

**Type** Package

**Title** Computes the p-values for Multinomial Goodness-of-fit tests

**Version** 1.0

**Author** Charles Thraves, Pablo Rivas, Felipe Subiabre

**Maintainer** Pablo Rivas <pablo.rivas.g@ug.uchile.cl>

**Description** Computes the p-value for an exact Multinomial, Pearson's Chi-Squared, log-likelihood ratio, or Power-Divergence tests. The main function, `pval_fourier`, uses a Fourier series expansion. Also, the library provides an exhaustive method `pval_exhaustive`, which evaluates all elements within the support of the Multinomial distribution. The function `pval_flexible` uses either of these methods, depending on the size of the instance.

**License** GPL ( $\geq 2$ )

**RoxygenNote** 7.3.2

**Encoding** UTF-8

## Contents

MultFourier-package . . . . .	1
pval_exhaustive . . . . .	2
pval_flexible . . . . .	4
pval_fourier . . . . .	6
<b>Index</b>	<b>8</b>

---

MultFourier-package	<i>MultFourier: Efficient Computation of Multinomial Test p-values</i>
---------------------	--

---

## Description

Computes the p-value for an exact Multinomial, Pearson's Chi-Squared, log-likelihood ratio, or Power-Divergence tests. The main function, `pval_fourier`, uses a Fourier series expansion. Also, the library provides an exhaustive method `pval_exhaustive`, which evaluates all elements within the support of the Multinomial distribution. The function `pval_flexible` uses either of these methods, depending on the size of the instance.

## Details

The MultFourier package provides efficient methods for computing p-values in multinomial tests, including both exact and approximate approaches. The exact p-value is computed using an exhaustive algorithm, which is recommended for small datasets. For larger datasets, an approximation is provided using a Fourier series algorithm, which is computationally efficient and provides control over precision through parameters such as `max_terms`, `rel_eps`, and `undersampling`.

The package includes the following methods:

- `pval_exhaustive`: Computes the exact p-value by enumerating all possible outcomes. Suitable for small datasets.
- `pval_fourier`: Approximates the p-value using a Fourier series expansion. Efficient for large datasets.
- `pval_flexible`: Automatically selects between the exhaustive and Fourier series methods based on the dataset size.

Key features:

- Control over precision through parameters such as `max_terms`, `rel_eps`, and `undersampling`.
- Support for multiple test statistics, including logP, Pearson's chi-square, and log-likelihood ratio.
- Verbose output for debugging and analysis.

## Author(s)

**Maintainer:** Pablo Rivas <pablo.rivas.g@ug.uchile.cl>

**Authors:**

- Charles Thraves
- Felipe Subiabre

## References

- Thraves, C. and Subiabre, F. et al. (Year): "Title of the Paper". Journal Name, Volume(Issue), Pages.

## See Also

- Report bugs at <https://github.com/PabRvss/MultFourier/issues>

---

pval\_exhaustive

*Compute p-value using an exhaustive method*

---

## Description

This function computes the p-value for an exact Multinomial, Pearson's Chi-Squared, log-likelihood ratio, or Power-Divergence tests. It evaluates all elements within the support of the Multinomial distribution to compute the exact p-value. Works for small-size instances.

**Usage**

```
pval_exhaustive(
  x,
  p,
  stat = 'prob',
  lambda = 0,
  max_time = 600,
  verbose = FALSE
)
```

**Arguments**

x	An Integer vector with realizations for each category.
p	A Numeric vector with the probabilities for each category. These should be non-negative and sum to one. It should be the same size as 'x'.
stat	String with the name of the statistic to compute. If 'prob', the exact Multinomial p-value is computed. If 'pearson', the Pearson's Chi-square p-value is computed. If 'llr', the log-likelihood ratio p-value is computed. If 'power_div', a Power Divergence p-value is computed, in which case a 'lambda' parameter must be given. The default value is "prob".
lambda	A Numeric with the lambda value of the Power Divergence statistic. Only works if stat = 'power_div', otherwise is ignored.
max_time	A Numeric with the maximum time limit in seconds. The default is 600.
verbose	Boolean. If 'TRUE', it prints information on the run time. If 'FALSE', it does not print. The default is 'FALSE'.

**Value**

Returns a 'S3' object with the following attributes:

- 'x': The input vector of the observed realizations for each category.
- 'p': The input vector of the probabilities for each category.
- 'pval': The p-value computed.
- 'time': The total execution time of the algorithm in seconds.
- 'p0': Probability mass function evaluated in 'x'.
- 'status': The final status ID of the algorithm upon completion:
  - '0': Successful computation.
  - '1': Maximum time reached.
- 'message': The finishing status displayed as a message.
- 'method': A String with value 'exhaustive'.

**Examples**

```
# Example 1: Compute p-value using the exhaustive method
probs <- c(0.1, 0.2, 0.3, 0.4)

x0 <- rep(10,4)

result <- pval_exhaustive(x0, probs, verbose = TRUE)
```

```

print(result)

# Example 2: Error case (mismatched lengths)
## Not run:
x0_invalid <- c(10, 10, 10)
probs_invalid <- c(0.2, 0.3) # Different lengths, should raise an error
result_invalid <- pval_exhaustive(x0_invalid, probs_invalid)
print(result_invalid)
## End(Not run)

```

---

pval\_flexible

---

*Compute p-value for a Multinomial Test*


---

## Description

This function computes the p-value for an exact Multinomial, Pearson's Chi-Squared, log-likelihood ratio, or Power-Divergence tests. It uses full enumeration if the cardinality of the support of the Multinomial distribution is less than  $10^6$ ; otherwise, it uses the Fourier method.

## Usage

```

pval_flexible(
  x,
  p,
  stat = "prob",
  lambda = 0,
  max_time = 600,
  max_terms = 300,
  rel_eps = 0.001,
  undersampling = 1,
  verbose = FALSE
)

```

## Arguments

x	An Integer vector with realizations for each category.
p	A Numeric vector with the probabilities for each category. These should be non-negative and sum to one. It should be the same size as 'x'.
stat	String with the name of the statistic to compute. If 'prob', the exact Multinomial p-value is computed. If 'pearson', the Pearson's Chi-square p-value is computed. If 'llr', the log-likelihood ratio p-value is computed. If 'power_div', a Power Divergence p-value is computed, in which case a 'lambda' parameter must be given. The default value is 'prob'.
lambda	A Numeric with the lambda value of the Power Divergence statistic. Only works if stat = 'power_div', otherwise is ignored.
max_time	A Numeric with the maximum time limit in seconds. The default is 600.
max_terms	An Integer indicating the number of terms to add in the Fourier series. The default is 300.
rel_eps	A Numeric with the relative error tolerance. The default is 0.001.

undersampling	An Integer with the undersampling value to use. The default and recommended value is 1. Values greater than one will speed up calculations but will sacrifice precision.
verbose	Boolean. If 'TRUE', it prints intermediate results every 10 terms. If 'FALSE', it does not print intermediate computations. The default is 'FALSE'.

## Value

Returns a 'S3' object with the following attributes:

- 'x': The input vector of the observed realizations for each category.
- 'p': The input vector of the probabilities for each category.
- 'pval': The p-value computed.
- 'gamma': The optimal gamma obtained in the first part of the method.
- 'n\_terms': The number of terms of the Fourier sum.
- 'time': The total execution time of the algorithm in seconds.
- 'p0': Probability mass function in 'x'.
- 'status': The final status ID of the algorithm upon completion:
  - '0': Converged.
  - '1': Maximum time reached.
  - '2': Maximum number of terms reached.
  - '3': Could not solve the optimization of gamma.
- 'message': The finishing status displayed as a message.
- 'method': A String with value 'fourier' or 'exhaustive', depending on the method used.

## Examples

```
# Example 1: Compute p-value using the exact multinomial statistic
probs <- c(0.00040161, 0.00080321, 0.00200803, 0.00401606, 0.00682731,
           0.01044177, 0.01485944, 0.02008032, 0.02610442, 0.03293173,
           0.04056225, 0.04899598, 0.05823293, 0.06827309, 0.07911647,
           0.09076305, 0.10321285, 0.11646586, 0.13052209, 0.14538153)

x0 <- rep(10, length(probs))

result <- pval_flexible(x0, probs, max_terms = 300, verbose = TRUE)
print(result)

# Example 2: Error case (mismatched lengths)
## Not run:
x0_invalid <- c(10, 10, 10)
probs_invalid <- c(0.2, 0.3) # Different lengths, should raise an error
result_invalid <- pval_flexible(x0_invalid, probs_invalid)
print(result_invalid)
## End(Not run)
```

pval\_fourier

*Compute p-value using Fourier method***Description**

This function computes the p-value for an exact Multinomial, Pearson's Chi-Squared, log-likelihood ratio, or Power-Divergence tests. It uses a Fourier series expansion. The infinite sum is truncated if any of the following happens: (i) runtime exceeds max\_time seconds, (ii) the modulus of o between the new sum term and the cumulative sum is less than rel\_eps for the last 40 iterations, or (iii) the number of terms of the sums is equal to max\_terms. The function allows to perform undersampling, but it is not recommended for every instance.

**Usage**

```
pval_fourier(
  x,
  p,
  stat = "prob",
  lambda = 0,
  max_time = 600,
  max_terms = 300,
  rel_eps = 0.001,
  undersampling = 1,
  verbose = FALSE
)
```

**Arguments**

x	An Integer vector with realizations for each category.
p	A Numeric vector with the probabilities for each category. These should be non-negative and sum to one. It should be the same size as 'x'.
stat	String with the name of the statistic to compute. If 'prob', the exact Multinomial p-value is computed. If 'pearson', the Pearson's Chi-square p-value is computed. If 'llr', the log-likelihood ratio p-value is computed. If 'power_div', a Power Divergence p-value is computed, in which case a 'lambda' parameter must be given. The default value is 'prob'.
lambda	A Numeric with the lambda value of the Power-Divergence statistic. Only works if stat = 'power_div', otherwise is ignored.
max_time	A Numeric with the maximum time limit in seconds. The default is 600.
max_terms	An Integer indicating the number of terms to add in the Fourier series. The default is 300.
rel_eps	A Numeric with the relative error tolerance. The default is 0.001.
undersampling	An Integer with the undersampling value to use. The default and recommended value is 1. Values greater than one will speed up calculations but will sacrifice precision.
verbose	Boolean. If 'TRUE', it prints intermediate results every 10 terms. If 'FALSE', it does not print intermediate computations. The default is 'FALSE'.

**Value**

Returns a 'S3' object with the following attributes:

- 'x': The input vector of the observed realizations for each category.
- 'p': The input vector of the probabilities for each category.
- 'pval': The p-value computed.
- 'gamma': The optimal gamma obtained in the first part of the method.
- 'n\_terms': The number of terms of the Fourier sum.
- 'time': The total execution time of the algorithm in seconds.
- 'p0': Probability mass function in 'x'.
- 'status': The final status ID of the algorithm upon completion:
  - '0': Converged.
  - '1': Maximum time reached.
  - '2': Maximum number of terms reached.
  - '3': Could not solve the optimization of gamma.
- 'message': The finishing status displayed as a message.
- 'method': A String with value 'fourier'.

**Examples**

```
# Example 1: Compute p-value using the Fourier method
probs <- c(0.00040161, 0.00080321, 0.00200803, 0.00401606, 0.00682731,
          0.01044177, 0.01485944, 0.02008032, 0.02610442, 0.03293173,
          0.04056225, 0.04899598, 0.05823293, 0.06827309, 0.07911647,
          0.09076305, 0.10321285, 0.11646586, 0.13052209, 0.14538153)

x0 <- rep(10, length(probs))

result <- pval_fourier(x0, probs, max_terms = 300, verbose = TRUE)
print(result)

# Example 2: Error case (mismatched lengths)
## Not run:
x0_invalid <- c(10, 10, 10)
probs_invalid <- c(0.2, 0.3) # Different lengths, should raise an error
result_invalid <- pval_fourier(x0_invalid, probs_invalid)
print(result_invalid)
## End(Not run)
```

# Index

## \* **package**

MultFourier-package, [1](#)

MultFourier (MultFourier-package), [1](#)

MultFourier-package, [1](#)

pval\_exhaustive, [2](#)

pval\_flexible, [4](#)

pval\_fourier, [6](#)